

long double	10 bytes	16 bytes
-------------	----------	----------

- **Enumerated type:** values are 0, 1, 2, ... by default; values can be set differently if desired
- **Variable declaration:** declarations reserve memory space (for local variables in the runtime stack); there are no default values: any uninitialized variables have **garbage values**
- **2 main scopes in C:** block scope and file scope; file scope variable are called **global variables**
- **Linkage:** controls whether different declarations refer to the same thing or different things
  - **No linkage:** each declaration refers to a different entity
  - **Internal linkage:** all occurrences of it inside a given file refer to the same entity, but occurrences in different files refer to different entities ( i.e., a different one in each source file)
  - **External linkage:** all occurrences of it refer to the same entity (i.e., there's only 1 in the whole program)
  - **By default:** identifiers with *file scope* have external linkage; identifiers with *block scope* have no linkage
  - **Changing linkage:**
    - Putting `extern` before a block scope identifier changes its linkage from none to external
    - Putting `static` before a file scope identifier changes its linkage from external to internal

### Lecture 3: Chp 3, 4, 5, 8

✓

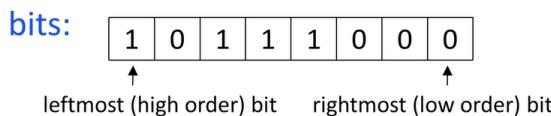
- **2 types of storage:**
  - **Automatic:** remains in memory only while it is in scope
  - **Static:** remains in memory the entire time the program is executing; global variables ALWAYS have static storage and cannot be changed
- **L-value and R-value:** left side of the equals sign is an L-value; right side of the equals sign is an R-value
- **Side effect:** an expression statement is only useful when the expression has a side effect
  - *Ex.* `y = 17;` or `x++;` or `z *= 3 + w++ - 4;`
  - *Not an example of a side effect:* `3 + w - 4;`
- **C has no boolean type:** 0 means “false” and any nonzero value means “true”
- **assert():** if its argument is “false” when a call to assert is executed, the program aborts with a message; use these for testing your code!
- **Symbolic constants:**
  - `#define` is another preprocessor directive
  - **Syntax:** `#define name value`
  - **Example:** `#define NUM_STUDENTS 543`
- \*Covers array basics; see **Chapter 8** for more information.
- \*Cover pass by value and pass by reference; see **Chapter 7 (Function Arguments)** for more information.
- \*Covers arithmetic operators; see **Chapter 5 (Types of Operators)** for more information.

## Lecture 4: Chp 5

✓

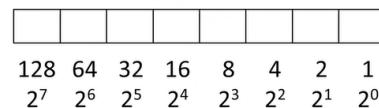
- **Conditional statements:**
  - Ex. `(x > y) ? x++ : y++;`  
If  $(x > y)$ , then  $x++$ . Else,  $y++$
- **L-value and R-value:** left side of the equals sign is an L-value; right side of the equals sign is an R-value
- **Promotion of char and short:** if a number value is stored in a `char` or `short`, they are promoted to an `int` before the expression is evaluated
- **Arithmetic conversions:** the “smaller” type is promoted to the “higher type in this order
- **Bit operations:**

`long double` ← highest type  
`double`  
`float`  
`unsigned long`  
`long`  
`unsigned int`  
`int`



- **Representation of unsigned integers:** in an unsigned number, bit position  $i$  represents the value  $2^i$  ( $i$  is 0 for the right most bit)

For an eight-bit unsigned integer:



- \*Covers bit shift operators; see **Chapter 5 (Bitwise & Bit Shifting)** for more information.
- \*Covers arithmetic operators; see **Chapter 5 (Types of Operators)** for more information.
- \*Covers equality, relational, and logical operators; see **Chapter 5 (Types of Operators)** for more information.

## Lecture 5: Chp 5

✓

- **Bit masking:** using the bit operators to apply a value to an integer that will isolate some of its bits to be able to examine or modify only those bits
  - See **Chapter 5 (Bitwise & Bit Shifting Chart)**
- \*Covers bit operators with binary numbers; see **Chapter 5 (Bitwise & Bit Shifting)** for more information.
- \*Covers Worksheet 1.

## Lecture 6: Chp 5, 6

✓

- **Mixed type assignments:** the value on the right side of an assignment (oritalization) is converted to the type of the L-value on the left before being stored (if their types are different)
  - For variables with values outside their type, the value wraps to fit
- **NULL pointer:** a pointer value that doesn’t point anywhere
  - It’s defined in a library header `stddef.h`, but can be used if you include any of the library header files

- **Pointers and pointer arithmetic:**
    - A function's return type can be a pointer
    - Pointers can point to array values
    - Pointers of the same type can be compared with == and !=
    - The operators ++ and -- can be applied to them
    - An integer value can be added or subtracted from a pointer
    - If 2 pointers are the same type, they can be compared with <, <=, >, and >=
  - **Casting example:** `int x= 10, *p= &x;`  
`double *q;`  
`q= (double *) p;`

\* and ++ are often combined with pointers  
(that are pointing to array elements)

– What do these do, and what is their difference?

**\*p++** What do we have to know in order to figure out what this means? **(\*p)++** **(\*p++)**

- If the prefix increment operator is used, the increment occurs first, however,  $*++p$  and  $++*p$  do not return the same value

Things work similarly for \* and --

- **const**: indicates a variable or parameter that can't be changed
  - **Typedef**: gives a new name to a type; Syntax: `typedef type-name new-name`
    - Ex. (note capitalization)      `typedef float Interest_rate;`
  - Most scalar types can be converted to other scalar types using **casts**
    - Ex.    `int *device = (int *) 100;`
  - **Generic (void) pointers**: can point to any type; can be assigned to any “real” pointer type variable (*and vice versa*) without casting    Ex. `void *p;`
  - \*Covers pointers and void pointers; see **Chapter 6** for more information.
  - \*Covers Worksheet 3.

## Lecture 7: Chp 7, 8

- \*Covers functions, prototypes, arguments, and parameters; see **Chapter 7** for more information.
  - \*Covers arrays vs. pointers; see **Chapter 8** for more information.
  - \*Covers multidimensional arrays and passing arrays to functions; see **Chapter 8** for more information.

## Lecture 8: Chp 8, 10

## Lecture 9: Chp 10, 11

✓

- **Unions:** looks like a structure but all its fields share the same memory space; only one of them “exists” or is “active” at a time
  - A common paradigm— include a union inside a structure, along with an enum that can be used to keep track of which union variant is currently active
  - A structure can have common fields and varying fields by having a union member
- **Stack memory vs. heap memory:** allocation and deallocation order
  - Why use dynamically-allocated memory:
    - To avoid compile-time limits
    - Linked data structures
  - Common errors:
    - Not checking the return value of `malloc()` or `calloc()`
    - Not allocating enough memory
    - Using `malloc()` and not initializing the memory returned if it needs to be initialized
  - Memory errors to avoid:
    - Calling `free()` on something not dynamically allocated
    - Freeing the same thing more than once (also undefined)
    - Calling `free()` on a pointer pointing anywhere other than the beginning of a dynamically allocated memory region

```
typedef struct {  
    int num_wheels;  
    double cost; /* how much $$$ */  
  
    enum { BICYCLE, CAR, BUS } vehicle_type;  
  
    union {  
        unsigned short num_speeds;  
        float miles_per_gallon;  
        unsigned int num_passengers;  
    } data;  
} Vehicle;
```

### Functions for allocating memory in `stdlib.h`:

`void *malloc(size_t size);`

- Allocates `size` bytes of memory from the heap (if enough is available) and returns a pointer to the beginning of it
- Does not initialize the memory at all!

`void *calloc(size_t num, size_t size);`

- Allocates enough memory (if enough is available) to store `num` things, each of size `size` bytes
- Returns a pointer to the beginning of the allocated memory, after initializing it to all 0 bits

Both return `NULL` if the requested amount of memory can't be allocated

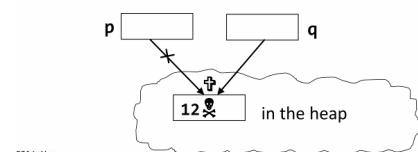
`void free(void *p);`

- Releases the memory pointed to by `p` and returns it to the free memory pool in the heap
- Always free memory when done with it!
- `free()` releases memory, but it does not:
  - change the values stored in the memory
  - change the value of the pointer it's called on
- Since `free()` does not change its parameter pointer, it's a good idea to set the argument to `NULL` right after the call (or point it somewhere else)

- `free(NULL)` is not an error (just does nothing)

- **Dangling pointer:** when an object is destroyed but the pointer's value has not been changed
- **Pointer aliasing:** See lecture video for in-depth explanation.

```
int *p= NULL, *q= NULL;  
p= malloc(sizeof(*p));  
q= p;  
free(p);  
p= NULL;  
... 0  
*q= 12;
```



## Lecture 10: EXAM WEEK

### Lecture 11: Chp 9, 11

- **realloc:** Check Chapter 11 (*Calloc and Realloc*)
  - void \*realloc(void \*ptr, size\_t new\_size);
- Recall that functions can return pointers: ----->
- You do not need to cast the return values of the memory allocation functions.

<b>BAD</b>	<b>GOOD</b>
int *p; p= <u>(int *)</u> malloc(sizeof(*p));	int *p; p= <u>malloc(sizeof(*p))</u> ;

- Garbage collection: less work for programmer; slower process; does not happen in C
- **NEVER** use **strtok** — Larry Herman
- Writing your own string function:

Suppose you wanted to write a string copy function, with prototype

```
void my_strcpy(char *q, const char *p);
```

You want to be able to call it like this:

```
char str1[80]= "", str2[80]= "duck";  
my_strcpy(str1, str2); /* copy str2 to str1 */
```

You really want to use pointer arithmetic  
(because you are an overachiever)

- Goes over Worksheet 7 (**Chapter 9**); Review **Chapter 9** and **Chapter 11**

✓

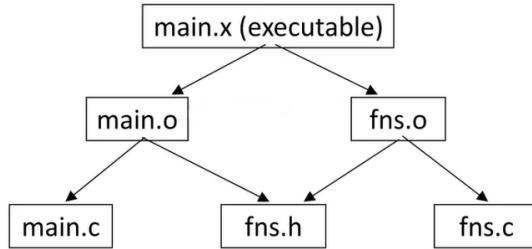
Recall that functions can return pointers

```
int *new_arr(int num, int initial_value) {  
    int *ptr= malloc(num * sizeof(int));  
    int i;  
  
    for (i= 0; i < num; i++)  
        ptr[i]= initial_value;  
  
    return ptr;  
}  
  
...  
  
int *values= new_arr(10, 42);
```

## Lecture 12: Make and Makefiles

✓

- **Compiling:**
  - You have to compile a source file every time you make a change, even if it's just changing 1 line. This is quicker to do for a small program; breaking a program up into smaller source files makes it more manageable.
- **Header file:** contains prototypes of functions so that they are defined before they are called
  - NEVER contains statements; includes enums, structure definitions, function prototypes, symbolic constants,typedefs, and sometimes extern global variable declarations
- **make:** utility that automates the process of building programs
  - Uses a **makefile** that specifies commands to create files and specifies which program files depend on which other ones
  - Anatomy of a **makefile**:



Composed of rules

```

target      dependency list
main.o: main.c fns.h
          gcc -c main.c ← action
          tab character

```

- A rule typically contains a *target*, a *dependency list*, and one or more *actions* (or directives)
- The target appears first (before a colon)
- The dependency list (after the colon) is the files that, if changed, should cause the target to be rebuilt
- Action lines follow a target, and must begin with a tab character

```

ss Makefile
main.x: main.o fns.o
        gcc main.o fns.o -o main.x

fns.o: fns.c fns.h
        gcc -c fns.c

main.o: main.c fns.h
        gcc -c main.c
grace10:~/<2>lectures/examples/lecture12/first-makefile-example:

```

```

main.o needs to be (re)created if main.c or fns.h changes
the command that will (re)create it gcc -c main.c

fns.o needs to be (re)created if fns.c or fns.h changes
the command that will (re)create it gcc -c fns.c

main.x needs to be (re)created if main.o or fns.o
the command that will (re)create it gcc main.o fns.o -o main.x

```

## • Running make:

### Running make

- See the complete makefile for the example on ELMS (just omitting some gcc options for now) in `first-makefile-example`
- If the command `make` is run it will read a makefile named `makefile` in the current directory, and try to build the first target in it
- If there isn't a makefile named `makefile` in the current directory but there's one named `Makefile` it'll be used instead
- The `-f makefile-name` make tells make to read a makefile with a different name

### Running make, con't.

- "make *target*" will try to build the target *target* instead of the first target in the makefile
  - You can also run make with multiple targets, e.g., "make *target1 target2 target3*"
- A target can have multiple actions, which will be executed sequentially (unless one of them has an error, in which case make will stop)
- The `-n` option lets you test a makefile by printing the commands that would be executed to build a target, without actually performing them (`make -n` or `make -n target`)

## • Modification time:

- Make decides what to do using files' modification times
- `ls` can display files' modification time:
  - `ls -l`
  - `ls -lt` modification time order
  - `ls -lrt` invert listing order
- What affects files' modification times?
  - Changing the contents of a file in any way
  - The `touch` command can be used (`touch file`) if you need to do it manually
    - Use only when really necessary

### Make's operation

- Make constructs a project dependency tree based on the makefile
- A target is considered to be out of date either if it doesn't exist, or if any one of its dependencies is newer than itself
- If a target is out of date, make performs its associated action
  - This process is performed bottom to top in the dependency tree

### Formal definition of make

- For any target *z*, note that *z*'s dependencies may be empty
- `make x` (*x* is a target) can be defined as follows:

```

make x:
for every d in x's dependencies where the makefile
contains a target named d
  perform "make d"
if ((x does not exist) or
  (x's dependencies has an element newer than x))
  perform x's action

```

- **Macros in makefiles:** a macro (makefile variable) is defined as a variable = value, and is used later by enclosing its name in `$(...)`
- For projects, you can only use features that are explained in this lecture. There will be **no leniency** for grading problems caused by using other features that you are being told not to use.

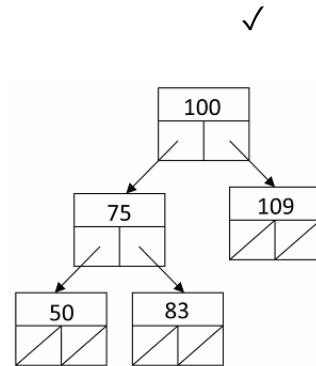
## Lecture 13: Chp 12, 13; Linked lists & Graphs

- Watch the lecture:

- **Linked data structures**
- **Singly-linked list**
- **Nodes**
- **Binary search tree** ----->

A typical declaration:

```
typedef struct node {
    int data;
    struct node *left;
    struct node *right;
} Node;
```



- **Directed graph:**

```

struct vertex; /* an incomplete type */

typedef struct edge {
    struct vertex *destination_vertex;
    struct edge *next_edge;
    /* perhaps other data for an edge here
       as well */
} Edge;

typedef struct vertex {
    Edge *edge_list;
    /* other data for a vertex here as well*/
} Vertex;
  
```

- **Doubly linked list:**

Inserting into a doubly-linked list

Suppose we want to create an ordered, doubly-linked list and maintain a tail pointer pointing to the list's last node at all times

```

typedef struct node {
    int data;
    struct node *next;
    struct node *prev;
} Node;
Node *head, *tail;
  
```

**Insertion must handle four cases:**

- inserting into an empty list,
- inserting a new first element,
- inserting a new last element, and
- inserting into the middle of a list

- **Advanced pointers:**

Inserting into a doubly-linked list

The four cases in more detail (try drawing pictures of these yourself):

– An initially empty list:

- Set the list's head and tail pointers (to a new node created)
- Set the new node's next and previous pointers to **NULL**

– Beginning of the list (add a new first element):

- Set the new node's previous pointer to **NULL**
- Set the new node's next pointer to the old first node
- Set the old first node's previous pointer to point to the new node
- Set the list head pointer to the new node

- How are declarations read in general in C?

```

int *p;
int *x[10];
int (*f)(void);
float (*z)(int *, char);
char *(*h[5])(int);

typedef void (*Ifp)(int);
Ifp p;

```

- See the examples `simple-fn-ptrs.c` and `fn-ptrs2.c`

- Review Chapter 13

## Lecture 14: Chp 13, 14

✓

- Command-line arguments:

`main()` can also be written with "parameters", as

```

int main(int argc, char *argv[]) {
    ...
}

```

In `main(int argc, char *argv[])`  
`argc` will have the number of arguments that appeared on the command line when the program was run, including the program's name  
`argv` will be an array of the arguments (as C strings) that appeared on the command line when the program was run  
`argv[0]` – the command that was invoked to start the program  
`argv[n]` – the  $n^{\text{th}}$  command line argument  
`argv[argc]` is always `NULL`

- Environment variables:

- `HOME`: contains the location of a user's home directory
- `PATH`: is used by the shell (and other programs) to find programs
- `envp`: will be an array of null-terminated strings, each of the form `KEYWORD = VALUE`, where `KEYWORD` is an environment variable name
  - Ex. `int main( int argc, char *argv[], char *envp[] )`

- Steps in compiling a C program:

- Source files are converted to object files
  - Preprocessing
  - Compilation (or translation of modified files):
    - Scanner
    - Parser
    - Type checker
    - Code generator
    - Optimizer
  - Compilers often produce assembly code, then run the assembler on it
- Object files are converted to an executable
  - This is done by the linker, which resolves symbols
    - Function use and definitions
    - Global variable declarations and external use

- **#define preprocessor directive**

- **Syntax:** `#define name replacement`
- `__FILE__`: the name of the source file being compiled
- `__LINE__`: the line number of the current line in the file
- `__DATE__`: the date the file was compiled
- `__TIME__`: the time the file was compiled
- `__STDC__`: 1 if the compiler conforms to ANSI C
- **Note:** these start and end with 2 underscores

- **Macros:**

- **Syntax:** `#define NAME (parameters) replacement`

- **Invocation:** `NAME (parameters)`

```
#define SUM(a, b) a + b
...
x= SUM(2, 5);  
x= SUM(6, 4);
```

- **Ex.**
- Macros are a single line, but you can define a "continuation" by ending a long line with a backslash:

```
#define COMPARE(s, t) (s.x == t.x ? \
0 : \
(s.x < t.x ? -1 : 1))
```

- Macros cannot be recursive
- The preprocessor doesn't check the types of macro arguments
- Macro substitution is not performed in string literals (or in comments)

- **Conditional Compilation:**

- `#if const-expr1`: code here is included by the preprocessor if and only if `const-expr1` is true
- `#elif const-expr2 /* optional */`: code here is included by the preprocessor if and only if `const-expr1` is false and `const-expr2` is true
- `#else const-expr2 /* optional */`: code here is included by the preprocessor if and only if `const-expr1` is false and `const-expr2` is false
- `#endif`

- One type of `const-expr` is `defined(NAME)`, which has the value 1 if `NAME` is a defined preprocessor symbol, and 0 otherwise
- The preprocessor directive `#error string` stops compilation at the point it appears, and prints `string`
- Conditional compilation is useful if code must be different on different systems, but overuse makes code hard to read

- **Nested or multiple inclusion of header files:**

- Usually done with all header files, so they can be included anywhere in any order without problems
- Use names of this format for conditional compilation— don't make up your own

A header file like `term.h` is always written as

```
#if !defined(TERM_H)
#define TERM_H

(all the previous contents of term.h appear here)

#endif
```

- File inclusion and makefiles:

What would dependencies in a makefile look like when one header file includes another one?

term.h:	main.c:
<pre>#if !defined(TERM_H) #define TERM_H definitions and prototypes here #endif</pre>	<pre>#include "polynomial.h" int main(void) {     code here }</pre>
polynomial.h:	
<pre>#if !defined(POLYNOMIAL_H) #define POLYNOMIAL_H #include "term.h" definitions and prototypes here #endif</pre>	

- Other preprocessor capabilities:

- `#undef NAME`
  - Undefines the preprocessor symbol `NAME`
- The compiler passes command-line definition of symbols to the preprocessor

```
gcc -D SIZE=100 file.c -o file.x
```

## Lecture 15: Chp 12



- Designing linked list operations:

Empty list case:

```
pre: curr == NULL == head == tail
post: new_node->prev == NULL == head == tail == curr
      new_node->next == NULL == head == tail == curr
      head == new_node
      tail == new_node
```

Beginning of list case:

```
pre: curr == head
post: new_node->prev == NULL == curr->prev
      new_node->next == head == curr
      head->prev == curr->prev == new_node
      head == new_node
```

Middle of list case:

End of list case:

```
pre: curr == NULL
post: new_node->prev == tail
      new_node->next == NULL == curr
      tail->next == new_node
      tail == new_node
```

```
pre: curr points to a node in the list
      data at curr is greater than value
      data at curr->prev is less than value
post: new_node->prev == curr->prev
      new_node->next == curr
      curr->prev->next == new_node
      curr->prev == new_node
```

- **Testing for errors:**

What do we try to test?

**Normal cases**

- Does the code work on valid inputs or parameters?

**Error cases**

- Does the code do what it's supposed to for invalid inputs or parameters, and what happens after an error (can we continue)?
- Often most bugs lie in this part of the code, since it has many infrequently-executed cases

**Environmental errors**

- Examples: out of memory, network is down, disk drive is full
- These situations can often be very hard to test

- **White box (or clear box) testing:**

- Create test cases by examining te statements (source code) of the code being tested
- Try to get maximum coverage of a set of tests

```

if (x)
    a;
if (y)
    b;
```

- Line or statement coverage

■ **Example set of tests:**

- $x = 1, y = 1$

- Branch coverage (at least one true, at least one false)

■ **Example set of tests:**

- $x = 1, y = 1$
- $x = 0, y = 0$

■ **Another example set of tests:**

- $x = 1, y = 0$
- $x = 0, y = 1$

- Path coverage:

■ **Example set of tests:**

- $x = 1, y = 1$
- $x = 1, y = 0$
- $x = 0, y = 1$
- $x = 0, y = 0$

- **Black box testing:**

- Create test cases based on the specifications of the code to be tested, without looking at the code
- Test for incorrect or missing functionality or behavior
- The problem with black-box testing is that edge, boundary, and corner cases may not be tested if tests are created without looking at the code

White-box testing- how to test different types of statements

- If-then-else- are all cases executed by at least some test?
- Switch- are all branches executed in at least some test?
- Loops
  - Is the loop run different numbers of time?
  - Is every possible termination condition tested by some test?
- Assignment statements- are a range of possible values assigned in different tests?

## Lecture 16: Chp 15, 16

✓

- **perror():**

In addition to `perror()`, the string library function `strerror(int err_number)` returns a (pointer to a) string describing an error, similar to what `perror()` prints (pass in `errno`)

- **Buffering:**

There are three types of buffering: *unbuffered*, *block buffering*, and *line buffering*

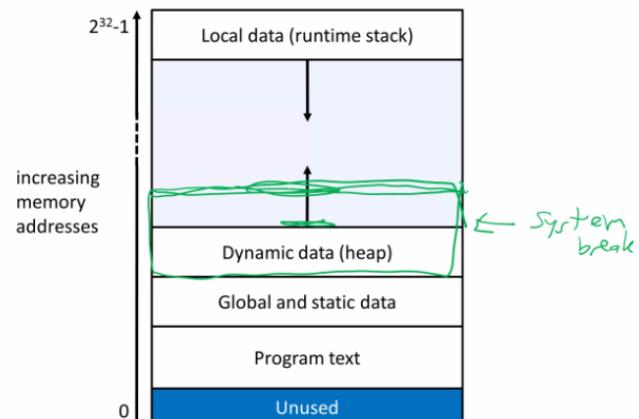
- Printing a newline also flushes a line-buffered stream

- **assert():**

- About `assert()` – if the preprocessor symbol `NDEBUG` is defined (including on the command line) when `<cassert.h>` was last included, any assertions don't have any effect at all
- In UNIX, a program using math library functions in `<math.h>` must be linked with the `-lm` option
- You can use the `man` command to view the manpage of any standard library functions

- **Implementing dynamic allocation:**

Virtual address space (from Lecture #1)



Where does memory for a program's heap come from?

- The memory allocation functions (`malloc()`, `calloc()`, and possibly `realloc()`) make calls to the OS (*system calls*) that add memory to the heap (move the top of heap pointer)
  - In Linux:  
`void *sbrk(int incr);`  
adds (at least) `incr` bytes to the program's data segment

For efficiency, the memory allocation functions get big chunks of memory from the OS (via `sbrk()`) when the heap size is increased, then hand out small regions to the program when allocation requests occur

How is the memory in the heap managed?

- Data structures, such as the *free list*, need to be maintained to keep track of what memory is in use and where the free memory regions are
- Calls to the memory allocation functions update these structures

- **Heap manager requirements:**

- The heap manager must be able to handle arbitrary sequences of requests to allocate and deallocate memory
- It must respond immediately
- All data it uses must be stored in the heap itself
- The heap must be *aligned* to be able to hold any type of data
- It can't modify already-allocated memory regions by increasing their size, changing their location, etc.- it will only modify free areas of memory
- \*Covers Worksheet 12.

#### Heap manager goals

- **Maximize throughput (speed)**
  - But it's often OK if the average memory allocation request is fast yet occasional calls are slow
- **Minimize fragmentation (wasted space)-**  
memory that is not available to satisfy allocation requests, even though it's not in use
  - There are two types of fragmentation:
    - *Internal fragmentation*
    - *External fragmentation*

## Lecture 17: Dynamic Memory Allocation

✓

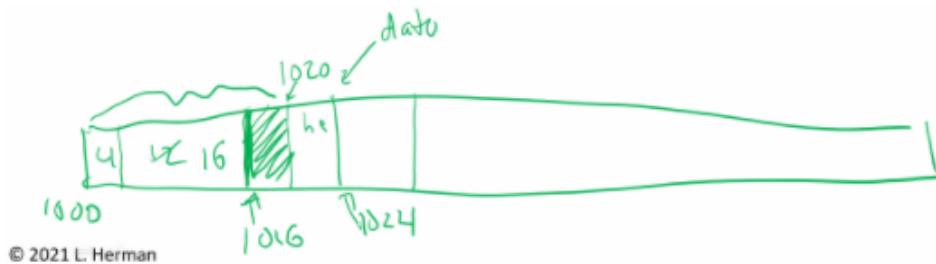
- **Heap:**

- The heap is divided up into blocks, which are formed of one or more 8-byte chunks
- Blocks will always begin at a memory address that is a multiple of 8
- Each block has an adjacent 32-bit header (just before it in the memory) that describes it



- **Block header format:**

- **29 bits** define the size of the block
- **1 bit** indicates whether the block is free or allocated
- **2 bites** are unused (some heap managers use them, but not here)
- The size stored in the header is the number of bytes allocated for data, plus the size of the header itself
- **4 bytes** of the memory before a block may be wasted, since the header must be located so that the following block begins on an address divisible by 8
- The end of the heap is indicated by a block with 0 in its size (its free/allocated bit is 1)



- **malloc() / calloc() using the simple heap- allocation; free() using the simple heap- allocation**

Using the simple heap- allocation

What the memory allocation functions

(`malloc()`/`calloc()`/possibly `realloc()`)

will do:

- They need to find a free block to satisfy the request:

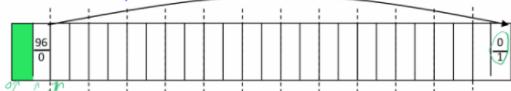
- They will start at the first block, and traverse the list of heap blocks, looking for a free one that's large enough to satisfy the request
- They'll mark the header of the selected block as allocated, and return a pointer to it (its memory address)

- **Initial heap:**

- Dash line is every 8 bytes
- Solid line is every 4 bytes

Example, assuming a small 104-byte heap

● **Initial heap:**



- The first number in a header is the size in bytes of the following block, including the header itself
- The second number in a header is the free/allocated bit – 0 means the space following that header is free and 1 means it's allocated (blue areas will be allocated memory)
- Green areas are wasted space, due to the necessity of aligning blocks on 8-byte boundaries

- **Two possible ways to choose (allocate) a block:**

- **First fit:** choose the first available block that's large enough
- **Best fit:** choose the smallest block that's large enough

- **Coalescing free blocks:**

- An adjacent sequence of free blocks can be merged into a larger free one
- A *coalescing* head implementation will update the header of the first of the adjacent free blocks to the size of the full free space

How can the prior block's header be found?

- Search from the beginning of the heap, or
- Include a footer at the end of each block, which also stores the block's size, or
- Include an explicit back pointer in each block's header, which stores the address of the previous block's header

Some allocation patterns may cause repeated splitting and merging

- Consider `p = malloc(n); free(p); q = malloc(n); free(q);` etc. If the first `malloc()` involves a split, the rest will too

Using the simple heap- deallocation

What `free()` will do:

- The block's header is four bytes before the pointer that `free()` is being called on – go to the header and mark the block as free

Review lecture 17 slides for the memory allocation example in detail.

What about splitting a large block into two blocks when allocating, one that's allocated, and one that's free?

- This gives the most efficient utilization of space at first, but it can waste space by leading to external fragmentation after repeated allocations and deallocations

- **Explicit free lists:**

An alternative to an implicit free list:  
explicit free lists

- **Maintain separate free lists of different sized chunks**

– For example, ones of 8, 16, 32, 64, 128, 256, 512, 1024, 4096 bytes, and anything larger than 4096

- **For an allocation request, return space from whichever list has chunks of the same size or just larger, if available**

- **Never split a block**

- **Automatically detecting heap problems:**

Automatically detecting heap problems

- **Auxiliary header storage – duplicate copies of headers can be stored elsewhere in the heap**

- **To enable detecting when programs write beyond the boundaries of allocated memory, a preamble and trailer can be added to each allocated block containing specific bit patterns**

- **If a free list of the needed size is empty:**

– Get one or more blocks from the next larger size free list, and if that fails, call `sbrk()` to get more storage from the operating system (some multiple of the size of the empty free list)

– Create a new free list of the needed size from the memory obtained

- **free():**

– Add the memory being freed back to the free list for the target size

– If all the chunks of a given size are free, release the space for that list to the common free pool

- **Virtual memory:**

- A program thinks it has memory addresses beginning with 0 (its virtual address space)
- Actual memory (the machine's *physical address space*) is divided up into pages (on Grace, 4K, or 4096 bytes)
- A program's memory can be loaded into physical memory anywhere
- When a program refers to a memory location, *address translation* is performed

- **Address translation:**

- The operating system maintains a data structure called they **page table** that keeps track of which physical pages a program's virtual pages are stored in
  - When a program refers to a memory location, the hardware and operating system (working together) split the memory address into 2 parts:
    - The low order bits (12 bits on a machine with 4K pages, because  $2^{12}$  is 4096) is an offset into a page
    - The system used the high order 20 bits as an index into the page table, which says where in the physical memory that page is

- **Simplified example: an instruction in a program refers to memory location 0x00200068.**

- The rightmost 12 bits are  $068_{16}$  ( $104_{10}$ )
- The high-order 20 bits are  $00200_{16}$  ( $512_{10}$ )
- Suppose the page table says that page 512 is stored in physical memory at actual starting address  $4000_{10}$
- Then the hardware and OS use actual memory location  $4104_{10}$  as the actual location that the instruction refers to

- **This all happens invisibly to the program**

- **This process is performed by the MMU (memory management unit) of the hardware**

- **Advantages of virtual memory**

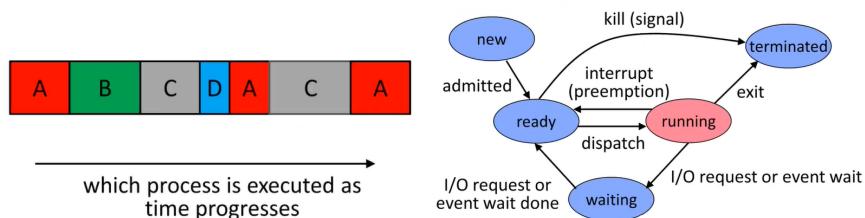
- A program's memory doesn't even have to be contiguous!
- A program can have a larger virtual address space than the machine's physical address space (i.e., a program can think it has more memory than the machine actually has)
- A program can be loaded anywhere into memory and run fine
- A program's virtual memory doesn't have to all be in physical memory at the same time
- Programs can easily share memory with the OS, and with each other, when needed

## Lecture 18: Process Control

✓

- Read the textbook reserve.
- The **kernel** is the core part of an operating system that is always in memory
  - It runs directly on the hardware and provides a more convenient abstraction of the machine for programs to use
- Modern systems have certain *privileged instructions* that can only be executed in **kernel mode**, not ordinary **user mode**
  - To execute these instructions, a normal user program has to do something to enter kernel mode, like making a **system call**
- **Multitasking:**
  - **Process:** writing the same code but with different data
    - **Program counter:** special register that always keeps the memory address of the next instruction to be fetched and executed
    - **Memory segment:**
      - Program code (often called the text segment)
      - Data (global variables and the heap)
      - Its runtime stack
    - **Contents of the processor registers:**
      - The program counter
      - General-purpose CPU registers
  - **Multitasking:** More than one process at the same time
    - **Data structures used to maintain the context of processes are:**
      - **The page table:** stores process' address spaces
      - **The process table:** stores information about each process
      - **The file table:** lists which files are opened by each process

Process state transitions



### ● Linux commands for viewing processes:

- **ps**: list current processes
  - **ps**, **ps -e**, and **ps -f**
- **pgrep name**: prints process number(s) of processes with *name* in their names
- **w**: shows who's logged in and what they're doing
- **top**
- **Ctrl-Z**, **bg**, **fg**, and **&**: for running a process in the background
- **jobs**: see the background processes currently running
- **pkill name**: kills processes with *name* in their names
- **kill process-ids**: kills processes given their process IDs

- **Signals:** a message to a process to notify it of an event

The kernel notifies processes of many events:

**SIGSEGV:** segmentation violation (a.k.a. segfault)

**SIGFPE:** floating point exception

**SIGILL:** illegal instruction

Users can trigger sending of signals:

**SIGINT:** interrupt (Ctrl-C)

**SIGQUIT:** quit and dump core (Ctrl-\)

**SIGTSTP:** terminal stop (Ctrl-Z)

- Processes have default actions when receiving signals (sometimes ignore, sometimes quit)
- There are mechanisms for processes to block signals, but two (**SIGKILL** and **SIGSTOP**) can't be blocked
- Processes can also send signals to each other (via the kernel)
- Signals are sent from the shell with the **kill** UNIX command, or by a program using the **kill()** system call
  - They do not always send the signal **SIGKILL**!

- **fork():**

The **fork()** system call

Creates a new copy of the calling process

```
#include <sys/types.h>
#include <unistd.h>
pid_t fork(void);
```

- It returns once in each process:
  - Once from the initial call
  - And also in the new process
- Its return value indicates which process is which:
  - 0 The child (new) process
  - > 0 The parent (original) process, and the value returned is the child's process ID
  - 1 The parent, but the child was not created (error case)

More things that are inherited by a child from its parent process:

- Process credentials: user and group ID (UIDs and GIDs in UNIX terminology)
- Open files (copies of *file descriptors*, to be covered shortly)
- Resource limits (that can be set and viewed with the tcsh **limit** command)
- Signal handling settings
- Process priority for scheduling by the operating system (called "nice" value in UNIX)

**fork()** semantics

Some things are inherited by a child from its parent process:

- A copy of the parent's memory contents, including program code, runtime stack, and heap
- The current working directory
- The environment
- The controlling terminal (the program that controls **stdin**, **stdout**, and **stderr** for the process, which is usually a shell), so the child reads input from and prints output to the same devices the parent does

Some things are unique to a child process:

- Its process ID
- Its parent process ID
- It has its own copies of its parent's file descriptors and directory streams (to be covered shortly)
- Its process times
- Its resource utilizations are initially set to 0
- Its pending signals are initialized to the empty set

- **Getting process IDs:**

Getting process IDs

```
pid_t getpid(void);
```

- Returns the process ID of the current process

```
pid_t getppid(void);
```

- Returns the process ID of the parent of the current process

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

- See the example **fork-example.c**

What if you overfork?

- The process table in the kernel can hold only a finite number of processes; what happens if it gets filled up?

• A **fork bomb:** `#include <unistd.h>`

```
int main(void) {
    while (1)
        fork();
    return 0;
}
```

- A fork bomb often requires sysadmin intervention (e.g., reboot, killing all user processes)

- Be careful when using **fork()** in a loop!

### Process termination

- A process can terminate itself, via the `exit()` system call
- The OS kernel can terminate a process
- One process can terminate another process, using the `kill()` system call
  - Can any process kill any other process?
    - No, that would be a Very Bad Thing™
    - For a process A to be able to kill another process B, either B has to belong to the same user as A, or A has to be a privileged process (i.e., kernel/root)

### Lecture 19: Chp

X

- Text

### Lecture 19: Chp

X

- Text

### Exam 1 - Last Minute Notes

#### Compilation Process

1. Pre-processing
2. Compilation
3. Assembly
4. Linkage

#### Scopes

- **File scope:** another way of saying “global variable”; always has static storage
  - `static`: putting this before a file scope identifier changes its linkage from internal to external
  - `Static` before a block variable changes its storage type but not its scope type
- **Block scope:** same as Java; coded in curly braces in a block
  - `Extern`: putting this before a block scope identifier changes its linkage from none to external

#### Static variables

- Can be used anywhere (and across functions) within the same source file

**Scalar types:** integer types, floating-point types, and pointer types; passed by value

**Void pointers:** cannot be dereferenced; can point to or be assigned to any type

**Side effect:** any modification to variables, files, etc. outside of a function’s local environment

**Typedef:** gives a new name to a type; Syntax: `typedef` type-name new-name

- Ex. (note capitalization)      `typedef float Interest_rate;`

#### Also know:

- Pass by value vs. Pass by reference

## **Exam 2 - Last Minute Notes**

---

- **String.h**
- **argc**: counts the number of arguments
- **argv**: the array that holds the arguments
- **Regression:**
  - **White box testing:**
  - **Black box testing:**
- **fgets():**
- **fscanf():**
- 

## **Exam 3 - Last Minute Notes**

---

Text